



[Course](#) > [Week 11](#) > [Final E...](#) > Q9: Rei...

## Q9: Reinforcement Learning

## Q9: Reinforcement Learning

Imagine an unknown game which has only two states  $\{A, B\}$  and in each state the agent has two actions to choose from:  $\{Up, Down\}$ . Suppose a game agent chooses actions according to some policy  $\pi$  and generates the following sequence of actions and rewards in the unknown game:

$t$	$s_t$	$a_t$	$s_{t+1}$	$r_t$
0	A	Down	B	2
1	B	Down	B	-4
2	B	Up	B	0
3	B	Up	A	3
4	A	Up	A	-1

Unless specified otherwise, assume a discount factor  $\gamma = 0.5$  and a learning rate  $\alpha = 0.5$

### Part 1

0.0/2.0 points (graded)

Recall the update function of Q-learning is:

$$Q(s_t, a_t) \leftarrow (1 - \alpha) Q(s_t, a_t) + \alpha (r_t + \gamma \max_{a'} Q(s_{t+1}, a'))$$

Assume that all Q-values initialized as 0. What are the following Q-values learned by running

Generating Speech Output above experience sequence?

$Q(A, Down)$

Answer: 1

$Q(B, Up)$

Answer: 7/4

### Explanation

Perform Q-learning update 4 times, once for each of the first 4 observations.

Submit

You have used 0 of 2 attempts

---

## Part 2

0.0/2.0 points (graded)

In model-based reinforcement learning, we first estimate the transition function  $T(s, a, s')$  and the reward function  $R(s, a, s')$ . Fill in the following estimates of T and R, estimated from the experience above. Write "x" if not applicable or undefined.

$\hat{T}(A, Up, A)$

Answer: 1

$\hat{T}(A, Up, B)$

Answer: 0

$\hat{T}(B, Up, A)$

Answer: 0.5

$\hat{T}(B, Up, B)$

Generating Speech Output

Answer: 0.5

$$\hat{R}(A, Up, A)$$

Answer: -1

$$\hat{R}(A, Up, B)$$

Answer: x

$$\hat{R}(B, Up, A)$$

Answer: 3

$$\hat{R}(B, Up, B)$$

Answer: 0

### Explanation

Count transitions above and calculate frequencies. Rewards are observed rewards.

You have used 0 of 2 attempts

To decouple this question from the previous one, assume we had **a different experience** and ended up with the following estimates of the transition and reward functions:

$s$	$a$	$s'$	$\hat{T}(s, a, s')$	$\hat{R}(s, a, s')$
A	Up	A	1	10
A	Down	A	0.5	2
A	Down	B	0.5	2
B	Up	A	1	-5
B	Down	B	1	8

ed on the above experience.

## Part 3

0.0/1.0 point (graded)

Give the optimal policy  $\hat{\pi}^*(s)$  for the MDP with transition function  $\hat{T}$  and reward function  $\hat{R}$ .

$\hat{\pi}^*(A)$

Up ▼

Answer: Up

$\hat{\pi}^*(B)$

Down ▼

Answer: Down

Submit

You have used 0 of 1 attempt

**i** Answers are displayed within the problem

## Part 4

0.0/1.0 point (graded)

Give  $\hat{V}^*(s)$  for the MDP with transition function  $\hat{T}$  and reward function  $\hat{R}$ .

Hint: for any  $x \in \mathbb{R}$ ,  $|x| < 1$ , we have  $1 + x + x^2 + x^3 + x^4 + \dots = 1/(1 - x)$ .

$\hat{V}^*(A)$

20

Answer: 20

$\hat{V}^*(B)$

16

Answer: 16

### Explanation

Based on the optimal policy in Part 3, calculate the value function using a Bellman equation.

Generating Speech Output

Submit

You have used 0 of 2 attempts

## Part 5

0.0/2.0 points (graded)

If we repeatedly feed this new experience sequence through our Q-learning algorithm, what values will it converge to? Assume the learning rate  $\alpha_t$  is properly chosen so that convergence is guaranteed.

☒ the values found above,  $\hat{V}^*$  ✓

☐ the optimal values,  $V^*$

☐ neither  $\hat{V}^*$  nor  $V^*$

☐ not enough information to determine

### Explanation

The Q-learning algorithm will not converge to the optimal values  $V^*$  for the MDP because the experience sequence and transition frequencies replayed are not necessarily representative of the underlying MDP. (For example, the true  $T(A, \text{Down}, A)$  might be equal to **0.75**, in which case, repeatedly feeding in the above experience would not provide an accurate sampling of the MDP.) However, for the MDP with transition function  $\hat{T}$  and reward function  $\hat{R}$ , replaying this experience repeatedly will result in Q-learning converging to its optimal values  $\hat{V}^*$ .

Submit

You have used 0 of 1 attempt

**i** Answers are displayed within the problem

© All Rights Reserved

Generating Speech Output